

The enterprise AI stack - a cheatsheet

Everything you need to know to have an informed conversation about AI in any enterprise setting. No jargon for the sake of jargon. If a concept is simple, it stays simple.

The 4 levels (and one you should probably ignore)

These are not steps you climb in order. They're a menu. Each business problem picks its own level. Most companies need level 1. Some need level 4. Almost none need level 5.

#	Level	What it does	When to use it	Cost (India)	Time
1	RAG	Connects an LLM to your company's data. System finds relevant documents and feeds them to the model before it answers. No training. Pure wiring.	Employees need to search internal docs, policies, contracts or knowledge bases using natural language.	10-30L to build, 2-5L/month to run	4-8 weeks
2	Fine-tuning	Changes how the model behaves - its tone, format, terminology. You feed it examples of how you want it to write. It learns your style.	Customer-facing content needs your brand voice. Legal docs need your terminology. Support replies need your tone.	5-15L per cycle	2-4 weeks
3	Distillation	Big model (teacher) answers thousands of questions. Small model (student) learns from those answers. 80% quality at 10% cost.	Paying too much for API calls. Need AI to run faster, cheaper or offline.	10-25L	4-6 weeks
4	Classical ML	Traditional ML on structured data. Random forests, gradient boosting, logistic regression. Predicts outcomes from patterns in numbers.	Structured data (spreadsheets, logs, transactions) and want to predict churn, risk, demand or failure.	20-50L to build, 5-10L/month	3-6 months
5	Own model	Train a base LLM from scratch. 99% of companies should ignore this entirely.	Only if your product IS the model. Otherwise, don't.	50-100Cr+	12-18 months

Services companies sell them level 5 ambition at 50Cr. The person who right-sizes the solution to the actual problem - that's who earns trust.

The agent layer - what sits on top

An agent is not a level. It's the orchestration layer that ties the levels together. You give it a goal in plain English. It breaks the goal into steps, picks the right level for each step, runs them in order and gives you one combined result.

Think of it like a project manager. You say "prepare me for tomorrow's meeting." The PM doesn't do the work himself. He looks at his team (RAG for search, fine-tuned model for writing, prediction engine for anticipation), assigns each task to the right person, collects the output and delivers one clean briefing.

The 5 agent concepts

Concept	What it means	Industry term
Task decomposition	Breaking a vague goal into concrete steps. "Prepare for meeting" becomes: predict files, search documents, write briefing.	Same everywhere. LangChain, CrewAI and AutoGen all use this term.
State management	Tracking what's done and what's next. A simple checklist: pending, running, done or failed.	Also called agent memory or scratchpad.
Tool selection	For each step, picking the right function to call. Search for finding, write for generating, predict for anticipating.	Called tool calling or function calling. When OpenAI or Anthropic say this, same concept.
Error recovery	When a tool fails, the agent doesn't stop. It logs the failure, skips the step and continues.	Called fault tolerance or graceful degradation.
Planning loop	Each step receives results from all previous steps as context. Step 3 knows what step 1 found.	Called ReAct (reason + act) in research papers.

The agent's quality depends on two things: how smart the planner model is (plan quality) and how good the tools are (execution quality). Better models make better plans. Better data makes better tools. The framework stays the same.

From agent to autonomous agent

A standard agent plans and executes. If a step produces garbage, it moves on without noticing. An autonomous agent adds reflection - after every step, it evaluates its own output. If the result is poor, it retries with a different approach. It catches its own mistakes.

The industry calls this the ReAct pattern (reason, act, observe, reason again). Four new concepts sit on top of the 5 agent concepts: **Evaluate** - after each step, check the result using simple rules (is it empty? too short? contains an error?) plus one LLM question (is this relevant?). **Retry** - when evaluation fails, change the approach. Simplify, rephrase or make it more specific. **Reflect** - log every attempt so you can see the agent's reasoning. **Converge** - max retries prevent infinite loops. Accept the best result and be honest about it.

Reflection can't fix a broken tool. But it tells you the tool is broken - instead of silently passing garbage to the next step. That's the difference between an agent that fails silently and an agent you can trust.

From autonomous agent to multi-agent

An autonomous agent talks to tools and evaluates itself. A multi-agent system has agents talking to each other. You give it a business decision. Two genuinely different models debate it from opposing perspectives. A judge reads the debate and synthesizes a verdict.

The key word is genuinely different. Most multi-agent demos use one model with different prompts - that's fake diversity. Using two different models (different architectures, different training data) produces genuinely different reasoning. Four new concepts: **Agent identity** - the LLM generates debate roles dynamically based on the topic. **Agent communication** - one agent's output becomes another agent's input. **Convergence** - after each round, agents

rate confidence (1-10). **Judge pattern** - a third call reads the full debate and synthesizes. Same pattern Microsoft ships as Critique in Copilot Researcher.

Two actually different brains is more diverse than a thousand copies of one brain with different names.

From agent to always-on agent

A standard agent waits for you to ask. An always-on agent runs on a schedule without being prompted. You define a goal ("find AI product jobs in India"), set a timer (7AM and 7PM daily) and let it run. It searches the internet, scores results against your context and stores them for review. The shift: from reactive (you ask, it answers) to proactive (it works while you sleep). This is where the model moves from a tool to a teammate.

The memory layer - what makes it an assistant

Without memory, every conversation starts from zero. You close the terminal, everything you said is gone. The documents are still there. But what YOU said, who YOU mentioned, what decisions were pending - vanished. A chatbot answers questions. An assistant remembers who you are.

The memory layer sits alongside your RAG. Every conversation, the system quietly extracts facts, entities and relationships. It stores them permanently. Next session, it retrieves relevant memories before the LLM sees your question. You never re-explain context.

Phase 1-8 gave your laptop intelligence about your FILES. Phase 9 gives your laptop intelligence about YOU. Your RAG knows what's in your documents. Your memory layer knows what's in your head - who you work with, what you're working on, what decisions are pending, what you care about. Two completely different knowledge sources, both feeding the same LLM.

The architecture

ChromaDB collection 1 (existing) - document chunks and image embeddings. This is your RAG. It knows what's in your files. **ChromaDB collection 2** (new) - memory facts extracted from every conversation. Each memory is an embedding with metadata: entity name, relationship type, timestamp, confidence score. **SQLite** (new) - entity relationship graph. Simple table: entity A, relationship, entity B, first seen, last seen, mention count.

The flow per turn

1. You ask a question. Python retrieves relevant memories from collection 2 and relationships from SQLite. These get injected into the system prompt alongside RAG chunks from collection 1. 2. The LLM generates an answer using all three sources: documents, memories and relationships. 3. Background extraction runs: the LLM reads the Q&A pair and extracts facts, entities and relationships as structured JSON. 4. You close the terminal. Documents stay in collection 1. Memories stay in collection 2. Relationships stay in SQLite. Tomorrow, the system already knows what you were working on.

Each turn uses roughly double the compute - one LLM call to answer, one to extract. On your laptop (Ollama), cost is zero. In the cloud, this is where Mem0's pitch lives: 1,800 tokens of distilled memories vs 26,000 tokens of raw history. 90% cheaper, 91% faster, 26% more accurate.

5 new memory concepts

Concept	What it means	Why it matters
Entity extraction	LLM identifies people, projects, companies and decisions from conversation text. Returns structured JSON. This is NER (named entity recognition) applied to conversations.	Without extraction, conversations are just text. With it, "Soham rejected the budget" becomes three linked facts.
Knowledge graph	SQLite stores entity-relationship-entity triples with timestamps and mention counts. Grows with every conversation. Connects facts across sessions.	You mention Soham in January, budget in February, "pending approvals" in March. The graph connects all three without you saying his name.
Memory consolidation	Handling contradictions and duplicates. "I live in Mumbai" in January, "I moved to Delhi" in March. System detects and updates. The hardest unsolved problem - Mem0 spent 18 months tuning this.	Without consolidation, memory becomes noise. With it, only the latest truth persists. Without a feedback loop, quality degrades over time.
Two-source retrieval	Every question triggers retrieval from both document RAG (collection 1) and memory (collection 2 + SQLite). Combined into one prompt.	Documents tell the LLM what's in your files. Memory tells it what's in your head. Both feed one answer.
Statefulness	LLMs are stateless by design. Every call starts from zero. Memory makes them stateful across sessions. A 128K context window can't survive a restart. Memory can.	On cloud, memory saves cost (1,800 tokens vs 26,000). On-device, memory saves context. Five precise facts beat 500 turns of raw chat.

Key concepts - what each word actually means

These build on each other. Start with how data gets into a model, then how models store knowledge, then how you compress and serve them.

Embeddings

A way to turn text into numbers so machines can compare meaning. "I'm hungry" and "I want food" look different as text but their embeddings are very close together. This is how RAG finds the right documents - it compares embeddings of your question against embeddings of every chunk in your database. Closest match wins.

Vector database

A database optimized for storing and searching embeddings. Regular databases search by exact match. Vector databases search by similarity (find chunks whose meaning is closest to this question). ChromaDB, Pinecone and Weaviate all do the same thing.

Chunking

Breaking large documents into smaller pieces before embedding them. You can't feed a 200-page PDF into an embedding model. So you split it into paragraphs (chunks), embed each separately and store them individually. When someone asks a question, the system retrieves the most relevant chunks.

Context window

The maximum amount of text an LLM can read at once. Think of it as the model's working memory. GPT-4 handles ~128K tokens (~100 pages). Mistral 7B handles ~8K tokens (~6 pages). Gemma 4: 128K tokens. If your chunks plus the question exceed the context window, the model can't read all of it.

Tokens

The unit LLMs think in. Not words - pieces of words. "Enterprise" is 2-3 tokens. "AI" is 1. Roughly: 1 token is about 3/4 of a word. 1,000 tokens is about 750

words. You pay per token. You hit limits in tokens. Everything is measured in tokens. Tokens are what go INTO training. What comes OUT of training are weights.

Weights and parameters

Every number stored in a model file is a parameter. Most are weights (control how much each input matters to the output) and some are biases (set the starting point when inputs are zero). An 8B model has 8 billion of these numbers. They start random and get nudged trillions of times during training - the model sees text, tries to predict the next token, gets corrected and adjusts. After trillions of corrections, those numbers encode the patterns from all the training data. The frozen numbers are the model file you download. When someone says "8B parameters" they mean 8 billion numbers. When someone says "open weights" they mean the company released the numbers but not the training data or code to reproduce the model.

Quantization

Compressing model weights to use less memory. Each weight is a number like 0.0217. At full precision (FP16), each weight takes 16 bits. Q8 = 8 bits (half the memory, almost no quality loss). Q4 = 4 bits (quarter the memory, some quality loss). A 7B model goes from 14GB at FP16 to ~4GB at Q4. This is how large models run on laptops. The naming system: Q4_K_M means 4-bit with k-quant (smarter grouping that gives important layers extra bits). Average BPW (bits per weight) for Q4_K_M is ~5, not flat 4. Below Q4, quality collapses - Q2 outputs incoherent garbage. Q4_K_M is the floor for production use.

Serving

The layer between a model file on disk and users who want answers. Without serving, a model is a file. With serving, it's an API endpoint applications can call over HTTP. Ollama does this on your laptop. vLLM does this at scale on GPUs. When you type in a terminal, you're the user. When code sends an HTTP request and parses JSON back, that's an application talking to a served model. Every chatbot, every API, every AI product has a serving layer.

Prefill vs decode

Every inference request has two phases. Prefill processes all input tokens in one parallel pass (fast). Decode generates output tokens one at a time - each depends on the previous one (slow). API providers charge less for input tokens (prefill is efficient) and more for output tokens (decode is 3-5x more expensive). The pause before Claude starts typing is prefill. The streaming text after is decode.

KV cache

During processing, the model computes a Key and Value for each token at each layer and stores them so it doesn't reprocess earlier tokens when generating the next one. KV cache grows linearly with conversation length. A 10,000-token conversation uses 100x more cache memory than a 100-token prompt. This is the largest memory cost in production serving. RAG applications with long retrieved context have large KV cache even on single-turn requests - it's total context that matters, not just conversation history.

TTFT and TPS

Two separate latency metrics. TTFT (time to first token) = how long until the first word appears (driven by prefill, grows with input length). TPS (tokens per second) = how fast output streams after that (driven by model size and hardware). Customer-facing chatbots optimize TTFT. Batch processing optimizes TPS. Different serving configs for each.

Throughput vs latency

Latency = time for one request to complete (what the user feels). Throughput = total tokens served per second across all users (what the system produces). Parallel processing can improve throughput but hurt latency per request. A CPO decides which to optimize based on the product.

Guardrails

Safety filters that wrap every model call. Input guardrails check the user's prompt before the model sees it (is it safe? on-topic?). Output guardrails check the model's response before the user sees it (contains PII? harmful content? wrong format?). Each check is a separate LLM call — guardrails triple the compute cost per turn. No single layer is reliable: input guardrails miss clever attacks, output guardrails catch most but not all. Defense in depth: regex for patterns (instant, free) + LLM for judgment (flexible, slow) + human review for edge cases (expensive, thorough).

Prompt injection

A user crafts their input to trick the model into ignoring its rules. "Ignore previous instructions" or "admin authorized override." The hardest unsolved problem in AI safety. General chat models used as safety classifiers catch nearly zero injection attempts. Purpose-built models like LlamaGuard exist for this reason. Anyone selling "100% safe AI" is lying.

Evaluation (evals)

A fixed set of test questions with expected answers, run against the model to measure quality. Not per-turn — run once before deployment and after any change. LLM-as-judge: a stronger model scores a weaker model's responses (1-10). If scores drop after a change, something broke. The eval set is an exam, not a textbook — never use eval questions for training. Easy eval questions give false confidence.

Data labeling

Tagging raw data with categories so a classifier can learn from it. Raw ticket "internet down for 3 days" gets labeled: category=network, urgency=high, sentiment=negative. Three methods: human labeling (accurate, expensive), regex rules (fast, brittle), LLM labeling (scalable, needs verification). Labeling organizes data — it doesn't change the model's weights. That's the difference from fine-tuning. Weak labeling model = bad labels = bad classifier = bad production routing.

Monitoring

Tracking model performance over time. Not a new capability — it's evaluation (Phase 14) running on a schedule with alerting. Eval scores drop = model or data degraded. Guardrail block rate spikes = new attack pattern. Latency trends up = hardware or model issue. An always-on agent (Phase 10) can run this.

LoRA (low-rank adaptation)

A technique that makes fine-tuning possible on normal hardware. Instead of retraining all 7 billion parameters (needs massive GPUs), LoRA freezes 99.9% and only trains a tiny adapter layer (0.1%). Nearly as good as full fine-tuning but runs on a laptop.

Temperature

Controls how creative vs predictable the model is. Temperature 0 = always picks the most likely next word (factual, repetitive). Temperature 1 = more randomness (creative, sometimes wrong). For RAG and enterprise use, keep it low (0.1-0.3).

Inference

When the model actually runs and produces an answer. Training teaches the model. Inference uses what it learned. API costs are inference costs. Latency complaints are inference speed complaints. "Scheduled inference" means running the model automatically at a set time - what an always-on agent does.

TF-IDF

Term frequency - inverse document frequency. Measures which words in a question actually matter. "Show me the budget" - "show" and "me" appear everywhere (low importance). "Budget" appears rarely (high importance). TF-IDF gives "budget" a high score.

Feature engineering

Turning raw messy data into clean numbers a ML model can learn from. Your query log has timestamps and text. Feature engineering turns "Monday 9:15 AM" into day_of_week=1, hour=9, is_morning=1. The better your features, the better your predictions. This is the real skill in ML - not the algorithm.

Multimodal

A model that can process more than just text. Gemma 4 reads text and images. Two types: embedding models (convert images to vectors for search) and generation models (look at images and answer questions about them). Different jobs.

Shared latent space

When a text embedder and an image embedder produce vectors in the same coordinate system. A text query "architecture diagram" can match against both text chunks and images. This is what makes multimodal retrieval work without separate search pipelines.

Per-layer embeddings (PLE)

Gemma 4's trick. Instead of one giant embedding table shared across all layers, each layer gets a smaller specialized one. 8B total params, 4B effective at inference. Different from MoE (mixture of experts) which routes tokens to different sub-networks. Parameter count is not capability. Architecture efficiency is the game now. A well-designed 4B model can outperform a brute-force 8B model.

Conversational multimodal RAG

Text RAG retrieves documents and feeds them to a language model. Multimodal RAG does the same for images, PDFs, screenshots - any file type. The model can see both retrieved text and retrieved images, then answer questions grounded in both. Conversational means it remembers what you asked before.

The pipeline: scan files, embed text chunks with a text embedder, embed images with a vision embedder (both producing vectors in the same latent space), store in one vector database. At query time, search both separately (text and images get reserved slots so text doesn't dominate), pass retrieved context to a multimodal LLM, generate an answer.

Three failure points to debug in order: (1) extraction - did the PDF text come out clean or garbled? Data engineering problem. (2) retrieval - did the right chunks surface? Embedding/chunking problem. (3) generation - was the answer good given what was retrieved? Prompt engineering problem. Always check retrieval before blaming the prompt.

Cross-modal similarity (text query matching images) always scores lower than text-to-text. In production, generate text captions for images during ingestion and search those instead. Solves 80% of the image retrieval problem.

Token math - the constraint that shapes everything

Every LLM call consumes tokens. 1 token is roughly 4 characters in English. 500 characters of text is about 125 tokens. Everything in the prompt counts: system instructions, retrieved chunks, conversation history, your question. The model's context window is the hard ceiling.

Turn	System	Context	History	Question	Total	4K model	128K model
1	200	750	0	50	1,000	OK	OK
5	200	750	1,500	50	2,500	Tight	OK
10	200	750	3,500	50	4,500	Broken	OK
20	200	750	7,000	50	8,000	Broken	OK

Enterprise AI cost formula: context length x concurrent users x tokens per turn x hours per day.
Context length is a cost multiplier, not a feature. Doubling context can mean 4x compute (attention scaling).

Serving math - the constraints that shape cost

Token math tells you what fits in the model's brain. Serving math tells you what fits in your hardware and budget. Both constrain every production deployment.

Model memory (GB) = parameters (B) x bits per weight / 8

A 70B model at Q4_K_M (~5 BPW): 70 x 5 / 8 = 43.75 GB. On an 80 GB A100 GPU, that leaves ~36 GB for KV cache and compute overhead. BPW comes from the quantization step or model card. Reference: Q4_0 ≈ 4.0, Q4_K_M ≈ 5.0, Q5_K_M ≈ 5.5, Q8_0 ≈ 8.0, FP16 = 16.0

KV cache per token = measure from server output: KV buffer size / context length

For phi3 3.8B: 0.375 MB per token (verified empirically by running llama.cpp server). For 70B models: roughly 2-3 MB per token. Rule of thumb: KV per token scales with model size.

Max concurrent users = (GPU memory - model size - overhead) / (KV per token x avg context per user)

Example: A100 (80 GB), 70B model at Q4 (~44 GB), 2 GB overhead. Free = 34 GB. KV per user at 4,000 token avg = ~10 GB. Max users = 34 / 10 = 3 concurrent users per GPU. One GPU, three users. This is why companies use smaller models or harder quantization - or buy more GPUs.

Cost per request = (input tokens x input price) + (output tokens x output price)

Output tokens cost 3-5x more because decode is sequential while prefill is parallel. For self-hosted: cost = GPU time x hourly rate. GPU compute is 70-80% of variable serving cost.

Serving stack selection

Stack	When to use	Key strength
Ollama	Local development, single user, experimentation	Easy setup and model management
llama.cpp server	Self-hosted, CPU or single GPU, 1-10 users	Low-level control, configurable parallel slots
vLLM	Production GPU serving at scale	Continuous batching, PagedAttention for smart KV cache management
TGI (HuggingFace)	Production, teams already using HuggingFace	Tight HF ecosystem integration
Triton (NVIDIA)	Enterprise, multi-model serving, complex pipelines	Multi-model orchestration, NVIDIA support contracts

Start with Ollama for prototyping. Move to vLLM when you have real users. Consider Triton only when you have multiple models and enterprise requirements. Quantization decides model size in memory. KV cache decides per-user memory cost. Serving orchestrates both within available infra. All three are interdependent - a CPO doesn't tune them individually.

The strategic layer - what the bible doesn't cover

Step zero: data readiness

Before you pick any level, assess: is the data even in a state where AI can use it? This is where 60% of enterprise AI projects die - not in model selection, but in data cleanup. A product leader who walks in and says "before we discuss RAG or fine-tuning, show me where your data lives and in what format" saves the company 6 months of wasted effort.

Questions to ask: Where does the data live? How many systems? What format? Who owns it? Is there PII that needs masking? How current is it? If nobody can answer these in 30 seconds, the data is not ready.

Data readiness is the real bottleneck. Evaluation is harder than building. Time to first value kills more pilots than bad AI. Retrieval quality matters more than context window size. Consumer AI fails from overtrust, enterprise from undertrust.

How to measure if AI is working

Level	The question	How to answer it
RAG	Are we retrieving the right documents?	Have 50 employees rate answers weekly. Track % rated useful.
Fine-tuning	Does the AI sound like us now?	Blind test - employees can't tell AI from a colleague.
Distillation	Is the small model as good as the big one?	Run same 1,000 questions through both. Compare side by side.
ML predictions	Are the predictions actually right?	Of 10 flagged as risky, how many actually failed? That's precision.
Agent	Does the agent complete the goal end to end?	Give it 20 goals. Track completion rate and useful results.
Autonomous	Does it catch its own mistakes?	Compare output with and without reflection. Count quality catches.
Multi-agent	Do different models produce different reasoning?	Run same topic with 1 model vs 2. Compare argument diversity and verdict quality.
Multimodal	Does it retrieve the right images?	Query for visual content. Debug extraction first, then retrieval, then generation.
Memory	Does it remember what matters across sessions?	Start a new session. Reference something from last week. Did it surface without re-explaining?
Always-on	Does it surface useful results without being asked?	Review a week of scheduled runs. What % of results were worth reading?
Quantization	Did compression break quality?	Same prompt at Q4 and Q8. Compare output coherence and factual accuracy.
Serving	Can it handle concurrent users at target latency?	Fire parallel requests. Measure TTFT and TPS under load vs targets.

Build vs buy vs partner

Approach	When it makes sense	When it doesn't
Build	AI is your core product or a key differentiator. You have engineering talent. You need full control over data, models and iteration speed.	You're building AI because it sounds impressive, not because the problem demands it.
Buy (SaaS)	Standard use case - customer support, document search, content generation. Someone already built a good product.	Your data is sensitive and can't leave your infrastructure. The vendor's model doesn't understand your domain.
Partner	You need custom AI but don't have the team. A services partner builds it, you own the output.	You plan to run AI at scale long-term. Knowledge walks out when the partner leaves.

Who to hire

The most common enterprise mistake: hiring the wrong role for the wrong level.

Role	What they do	Covers	India salary
Software engineer	Builds the RAG pipeline, agent layer, autonomous agent framework, multi-agent systems, multimodal retrieval, memory layer and always-on agent infrastructure. APIs, data ingestion, vector DB wiring, reflection loops, agent communication, entity extraction, scheduled jobs.	Level 1 (RAG) + agents + autonomous + multi-agent + multimodal + memory + always-on	12-35L/year

ML engineer	Runs fine-tuning jobs, manages training pipelines, handles model deployment, benchmarks embedding models, handles multimodal alignment. Owns quantization decisions (which quant level for which use case) and serving infrastructure (vLLM, TGI setup, GPU memory budgeting, KV cache tuning).	Level 2 + 3 (fine-tuning, distillation) + quantization + serving + model selection	20-50L/year
Data scientist	Builds classical ML models from structured data. Designs features, tests hypotheses, evaluates predictions.	Level 4 (classical ML)	18-45L/year
Data engineer	Cleans data, builds pipelines, manages extraction from PDFs and documents, handles image captioning for multimodal search. The person who fixes step zero.	Always. Before anything else.	15-40L/year
Product manager	Decides what to build and why. Translates business problems into AI requirements. Measures outcomes. Owns quantization-quality tradeoff decisions and serving latency targets. Sets TTFT and TPS budgets.	Day one. Every phase.	25-60L/year
AI researcher	Designs new architectures, runs pre-training, publishes papers. Don't hire unless you're building a model company.	Level 5 only	40L-1Cr+/year

Common mistake: companies hire data scientists when they need software engineers (for RAG) or hire AI researchers when they need ML engineers (for fine-tuning). Wrong hire = 6 months wasted. Another: expecting software engineers to own serving optimization - that's ML engineering territory.

Red flags - what to watch for

These are the signals that an AI project, vendor or internal team is headed for trouble.

"We need our own model." No you don't. Unless your product IS the model, you need someone else's model plus your data. RAG + fine-tuning covers 95% of enterprise use cases.

"We'll start with a proof of concept." POCs without clear success criteria are where budgets go to die. Define what "working" looks like, who will judge it and what happens if it works. If the last answer is vague, the POC is theater.

"The data is ready." It almost never is. Where does it live? How many systems? What format? Who owns it? If nobody can answer in 30 seconds, the data is not ready.

"We need an AI strategy." You need a business problem that AI can solve. "Our support team takes 48 hours to resolve tickets - can AI bring that to 4?" is a real question. "We need an AI strategy" is not.

"We need GPUs." For RAG? No. For fine-tuning with LoRA? Your laptop works. For inference? API calls. For serving at scale? Then yes - but do the serving math first. How many concurrent users? What model size? What quant level? What latency target? If you can't answer these, you don't need GPUs yet - you need clarity.

"Let's hire a Chief AI Officer." A title won't save you. If you don't have a clear business problem, clean data and engineering capacity, a CAIO will spend 6 months writing PowerPoint decks. Level 1 needs a software engineer. Not a CxO.

"We need an agent." Most companies asking for agents haven't finished level 1. An agent orchestrates tools - if the tools don't work well, the agent just orchestrates garbage faster. Get RAG right first. Then fine-tune. Then add the agent layer.

"The AI hallucinated." It didn't hallucinate. Your retrieval served the wrong context. Check retrieval before blaming the model.

"More context window is always better." Context length is a cost multiplier. Doubling context can mean 4x compute due to attention scaling. And every extra token inflates the KV cache, which eats GPU memory that could serve more users. Right-size the window to the actual task.

"We'll just quantize to Q2 and save money." Below Q4, quality collapses. Q2 outputs incoherent garbage - tested and verified. The cost savings are meaningless when the output is unusable. Q4_K_M is the floor for production.

Quick reference - one-liners for each level

Designed to be said out loud, not read from a slide.

RAG: "Don't retrain the model. Just show it the right documents at the right time."

Fine-tuning: "RAG gives it knowledge. Fine-tuning gives it personality."

Distillation: "Make the expensive model teach the cheap model. Then fire the expensive model."

Classical ML: "LLMs are great at language. They're terrible at predicting numbers. Use the right tool."

Agent: "An agent without good tools is a project manager with no team. Fix the tools first."

Autonomous agent: "The agent that tells you 'this result is garbage' is more valuable than the one that says 'done.'"

Multi-agent: "Two different brains arguing is more honest than one brain pretending to be two."

Multimodal RAG: "Same retrieval pipeline, but now it can see images too. Debug extraction, then retrieval, then generation - in that order."

Memory: "A chatbot answers questions. An assistant remembers who you are."

Always-on agent: "An agent that doesn't wait for you to ask. It runs on a schedule and brings you answers."

Quantization: "Same 8 billion numbers, stored with less precision. 4x smaller, barely any quality loss."

Serving: "A model without serving is a file. A model with serving is a product."

Guardrails: "The prompt injection that beats both layers is the lesson. No single filter is enough."

Evaluation: "Evals are the exam, not the textbook. If you train on your test questions, you're cheating."

Data labeling: "Weak model labeling your data = garbage labels = garbage classifier. Spend on the labeler."

Monitoring: "Monitoring is not new. It's phases 12-15 running on repeat with an alert when numbers drop."

Own model: "If you're asking whether you need this, you don't."

The 16 phases

#	Phase	Concept	Model	Key learning
---	-------	---------	-------	--------------

1	RAG	Search by meaning	nomic-embed-text + Mistral 7B	Embeddings, vector search, chunking, retrieval + generation are separate steps
2	LoRA fine-tuning	Style transfer	TinyLlama 1.1B	Training pairs, adapter weights, style transfer
3	Distillation	Teacher-student	Mistral 7B teaching TinyLlama	Teacher-student, cost reduction, quality tradeoff
4	ML prediction	Classical ML	Random forest, scikit-learn	Feature engineering, cross-validation, scheduled inference
5	Standard agent	Goal to execution	Mistral 7B	Tools, planning, state, routing
6	Autonomous agent	Self-correction	Mistral 7B	ReAct pattern, reflection, self-correction
7	Multi-agent debate	Debate + judge	Mistral 7B vs Phi3 3.8B	Agent communication, convergence, judge pattern
8	Multimodal RAG	Text + vision	nomic-embed-vision + Gemma 4 E4B-IT	Vision embeddings, shared latent space, cross-modal retrieval, conversational context
9	Memory layer	Persistent context	Gemma 4 E4B-IT	Entity extraction, knowledge graph, memory consolidation, two-source retrieval, statefulness
10	Always-on agent	Scheduled search	Claude Haiku API + OpenClaw + headless browser	Scheduled inference, API-based model calls, headless browsing, runs 24x7 without prompting
11	Quantization	Model compression	Phi3 3.8B (FP16 → Q4_K_M via llama.cpp)	Weights/parameters/biases, BPW, Q2/Q4/Q8 quality-speed-size tradeoff, quantized a model from scratch
12	Serving	Model to API	Phi3 3.8B + Ollama API + llama.cpp server	Prefill vs decode, KV cache, concurrent request handling, GPU memory economics, serving stack selection
13	Guardrails	Safety filters	Phi3 3.8B (input + output classifier)	Input/output checks, prompt injection (0/5 caught at input, 4/5 at output), defense in depth, 3x cost per turn
14	Evaluation	Quality scoring	Phi3 3.8B (tested) + Gemma 4 E4B-IT (judge)	Eval sets, LLM-as-judge, regression testing, exam vs textbook, proved quantization didn't break quality
15	Data labeling	Tag raw data	Phi3 3.8B (LLM-as-labeler)	Auto-labeling, format failures (30%), weak model = bad labels, labeling ≠ fine-tuning
16	Monitoring	Track over time	Conceptual (combines phases 12-15)	Drift detection, alerting, scheduled evals, phases 12-15 on repeat

All 16 phases built on a 50K laptop · no GPU · no cloud · no API key (except phase 10) · open-sourced on GitHub
Part of the Laptop AI project - github.com/swapnil-shah-ai/laptop-ai
Built by learning, not by reading. Every concept on this page was implemented on a regular laptop with no GPU.